

How Good Are GPT Models at Machine Translation? A Comprehensive Evaluation

Amr Hendy, Mohamed Abdelrehim, Amr Sharaf, Vikas Raunak,
Mohamed Gabr, Hitokazu Matsushita, Young Jin Kim,
Mohamed Afify, Hany Hassan Awadalla*
Microsoft

Abstract

Generative Pre-trained Transformer (GPT) models have shown remarkable capabilities for natural language generation, but their performance for machine translation has not been thoroughly investigated. In this paper, we present a comprehensive evaluation of GPT models for machine translation, covering various aspects such as quality of different GPT models in comparison with state-of-the-art research and commercial systems, effect of prompting strategies, robustness towards domain shifts and document-level translation. We experiment with eighteen different translation directions involving high and low resource languages, as well as non English-centric translations, and evaluate the performance of three GPT models: ChatGPT, GPT3.5 (text-davinci-003), and text-davinci-002. Our results show that GPT models achieve very competitive translation quality for high resource languages, while having limited capabilities for low resource languages. We also show that hybrid approaches, which combine GPT models with other translation systems, can further enhance the translation quality. We perform comprehensive analysis and human evaluation to further understand the characteristics of GPT translations. We hope that our paper provides valuable insights for researchers and practitioners in the field and helps to better understand the potential and limitations of GPT models for translation.

1 Introduction

Recent advancements in natural language processing (NLP), particularly the development of large-scale language modeling techniques, have brought remarkable improvements in machine translation as well as in other NLP tasks (Fan et al., 2021; Kim et al., 2021; Costa-jussà et al., 2022). The emergence of large language models with diverse capabilities, including machine translation, has opened

up new possibilities for building more effective translation systems (Brown et al., 2020; Chowdhery et al., 2022). Among these models, the latest Generative Pre-trained Transformer (GPT) models (Brown et al., 2020) have gained significant attention for their ability to generate coherent and context-aware text. We present a comprehensive evaluation of GPT models for machine translation, exploring their strengths and limitations, and providing insights for researchers and practitioners working in the area of machine translation.

GPT models and the conventional Neural Machine Translation (NMT) systems are both based on the transformer architecture (Vaswani et al., 2017), but they differ in several aspects. First, GPT models are decoder-only models that use the same parameters to process the context and the source as a single input for generating the next output. On the other hand, NMT models usually have an encoder-decoder architecture that encodes the source sentence in the encoder network and decodes the target sentence conditioned on the previous outputs in the decoder network. Second, GPT models are mainly trained on monolingual data, with a strong bias towards English¹, whereas NMT models rely on large amounts of highly curated parallel data. Third, GPT models need a much larger number of parameters to achieve multilingual in-context capabilities. We have observed that GPT models exhibit promising translation capabilities, even with these differences in architecture and training data.

The performance of GPT models in machine translation, despite their promising potential, remains under-investigated relative to commercial and state-of-the-art research systems. This study aims to address this research gap by systematically assessing the efficacy of GPT models for machine translation, with a focus on their performance,

Corresponding author: hanyh@microsoft.com

¹https://github.com/openai/gpt-3/blob/master/dataset_statistics/languages_by_character_count.csv

prompts, document-level translation, domain robustness and possible advantages of integrating them with conventional NMT systems.

To explore the potential of GPT models for translation, we perform comprehensive experiments to examine their translation abilities. Specifically, we investigate the performance of GPT models for machine translation across 18 language pairs, covering both high and low-resource languages, as well as English-centric and non-English-centric directions. We compare the quality of three GPT models: *text-davinci-002*, *text-davinci-003* (GPT3.5), and *ChatGPT*, and show that they differ significantly in their translation capabilities.

We also explore the impact of prompting strategies on the performance of GPT models for machine translation. We examine both the content and the form of the prompts, and identify the best practices for obtaining optimal results. Furthermore, we test the hypothesis that GPT models would enhance document-level translation, as they could exploit the context and coherence of the entire document to generate more accurate and fluent translations. We evaluate this hypothesis on various language pairs using several metrics. Additionally, we evaluate the cross-domain generalization ability of GPT models for translation tasks and examine their robustness under domain shift.

Moreover, we conduct extensive human evaluations and analyses to provide valuable insights into the strengths and weaknesses of GPT models for machine translation, and to suggest directions for future work. We also perform comprehensive analyses to understand whether GPT and NMT models have complementary characteristics, and we propose several ideas to combine the advantages of the two paradigms. Finally, we touch on the effectiveness of GPT models on cross-lingual natural language tasks beyond translation, and explore their multilingual capabilities and limitations.

To explore the research questions above, we organize the paper as follows:

- We provide a detailed experimental setup (§2), which includes the datasets used (§2.1), the machine translation systems used in the comparison (§2.2), the GPT systems (§2.3), and the evaluation methods (§2.4).
- We present a series of experiments (§3) that investigate different aspects of GPT models for machine translation. These experiments cover

prompt selection strategies (§3.1), zero-shot translation capabilities of GPT models (§3.2), GPT performance on high-resource languages (§3.3), GPT performance on low-resource and non-English-centric languages (§3.4), document-level MT with GPT (§3.5), translation robustness toward domain shift (§3.6), and hybrid GPT and NMT translation (§3.7).

- We then present the human evaluation and analysis that provides insights into the quality of GPT translations (§4).
- We discuss the characteristics of GPT translations (§5) vis-à-vis NMT and analyze the differentiating aspects of GPT translations (§5.1) by quantitatively enumerating language modeling bias artifacts (§5.2), the characteristics of translation across various language directions (§5.3-§5.5), as well as parallel data bias artifacts (§5.6).
- We explore the multilingual capabilities of GPT models beyond translation (§6).
- We conclude by summarizing our findings and suggesting future directions for research (§7).

2 Experimental Setup

2.1 Datasets

We considered 18 different translation directions across a diverse set of languages for our comprehensive evaluation. The evaluation covers both high- and low-resource languages, as well as English-centric and non-English-centric direct translations. The languages considered in this study include European (English-*EN*, French-*FR*, German-*DE*, Czech-*CS* and Icelandic-*IS*), Asian (Chinese-*ZH* and Japanese-*JA*), Cyrillic (Russian-*RU* and Ukrainian-*UK*) and African (Hausa-*HA*).

We use publicly available datasets to facilitate reproducibility and data sharing. We use the WMT22 testsets² for all languages except for Icelandic and Hausa, for which we use the WMT21 testsets³. We use WMT22 datasets for two reasons: first, they are recent and less likely to overlap with the GPT models training data that was collected until June

²<https://www.statmt.org/wmt22/translation-task.html>

³<https://www.statmt.org/wmt21/translation-task.html>

2021⁴. Second, they have natural source texts and translated target texts, which avoid the problems of using “translationese” testsets with unnatural source texts in the original language that may result in inaccurate evaluation (Zhang and Toral, 2019).

Table 1 summarizes the datasets and sizes used in this paper. We focus on the most recent datasets with directional translation, as older datasets may have influenced the training data of the GPT models. We make all data and analysis in this paper publicly available to promote further research.⁵

Lang-Pair	Dataset	Number of sentences	WMT-Best System
CS-EN	WMT22	1448	Online-W
EN-CS		2037	Online-W
DE-EN	WMT22	1984	Lan-Bridge
EN-DE		2037	Online-B
IS-EN	WMT21	1000	Facebook-AI
EN-IS		1000	Facebook-AI
JA-EN	WMT22	2008	DLUT
EN-JA		2037	NT5
ZH-EN	WMT22	1875	JDExploreAcademy
EN-ZH		2037	Online-W
UK-EN	WMT22	2018	Lan-Bridge
EN-UK		2037	Online-B
RU-EN	WMT22	2016	JDExploreAcademy
EN-RU		2037	Online-W
HA-EN	WMT21	997	Facebook-AI
EN-HA		1000	Facebook-AI
FR-DE	WMT22	2006	Online-W
DE-FR		1984	Online-B

Table 1: Test datasets used in the evaluation and best systems used for comparison as reported in WMT by Kocmi et al. (2022b) and Barrault et al. (2021).

2.2 Neural Machine Translation Systems

In this study, we compare the performance of GPT systems against both state-of-the-art (SoTA) research and commercial systems. We use the top ranked systems (WMT-Best) in the WMT evaluation campaigns for each language pair, which we use as a baseline for comparison. WMT-Best systems are a mix of top ranked commercial and research systems. We use the system outputs as provided by the evaluation campaigns (Kocmi et al., 2022b; Barrault et al., 2021). Table 1 shows the list of top ranked system for each language pair. We also utilize Microsoft Translator which we access through the public API available on Azure

Cognitive Services⁶.

2.3 GPT Systems

We assess the latest three variants of the largest GPT models available, as listed at OpenAI’s documentation⁷. These models are:

- *text-davinci-002* - an InstructGPT model (Ouyang et al., 2022) which utilizes Reinforcement Learning with reward models trained based on human comparisons.
- *text-davinci-003* - an improved version of *text-davinci-002*.
- *ChatGPT* - a model that is similar to the previous two and optimized specifically for conversational purposes⁸.

All GPT models have been accessed through APIs on Microsoft Azure OpenAI Service⁹.

2.4 Evaluation Methods

Sentence-Level Evaluation The MT Metrics shared task (Freitag et al., 2022) recommends the use of neural network-based metrics in machine translation evaluation, as they have demonstrated a high correlation with human evaluation and are resilient to domain shift. Following these recommendations, we employ the top-ranked metrics from *Unbabel* in the shared task¹⁰. Specifically, we use COMET-22 (*wmt22-COMET-da*) (Rei et al., 2022a), a reference-based metric that combines direct assessments (DA), sentence-level scores, and word-level tags from Multidimensional Quality Metrics (MQM) error annotations. For referenceless quality estimation, we adopt COMETkiwi (*wmt22-COMETkiwi-da*) (Rei et al., 2022b). Additionally, we report results using *SacreBLEU*¹¹ and *Chrf* (Popović, 2015) for completeness, although we note that these metrics are not extensively used in our analysis.

Document-Level Evaluation For the experiments on document-level translation using GPT, we face a challenge in evaluating performance due

⁴<https://help.openai.com/en/articles/6639781-do-the-openai-api-models-have-knowledge-of-current-events>

⁵<https://github.com/microsoft/gpt-MT>

⁶<https://azure.microsoft.com/en-us/products/cognitive-services/translator>

⁷<https://beta.openai.com/docs/model-index-for-researchers>

⁸<https://openai.com/blog/chatgpt>

⁹<https://azure.microsoft.com/en-us/products/cognitive-services/openai-service>

¹⁰<https://github.com/Unbabel/COMET>

¹¹<https://github.com/mjpost/sacrebleu>

to the lack of metrics that can handle the non one-to-one sentence mapping that the systems may produce. To address this challenge, we have adapted the COMET metrics to better suit document-level evaluation. The adaptation involves splitting the document into multiple segments with an overlapping sliding window and computing the average score across these segments to compare two documents. We use Doc-COMETkiwi to refer to the modified metric throughout the evaluation.

This simple modification has three clear design advantages over a pure sentence-level evaluation. First, it allows each sentence to be evaluated within its context. Second, it enables each sentence to be evaluated across multiple contexts due to the overlapping nature of the sliding window. Lastly, it avoids the limited context window of the evaluation models that could hinder quality assessment over longer static windows on documents. We do not claim that this is an optimal metric for evaluating document-level translation, but it overcomes the limitation of the one-to-one sentence mapping and may capture the quality of translation in ambiguous contexts better than sentence-level representation. We argue that developing more robust document-level metrics is still essential.

The current metrics for evaluating machine translation performance may not be adequate for measuring the performance of GPT models, and it may be necessary to develop new metrics that take into account the unique characteristics of these models.

Human Evaluation and Analysis We perform human evaluation (§4) using source-based sentence-level contrastive Direct Assessment + Scalar Quality Metric (contrastive DA+SQM; Akhbardeh et al. 2021, Kocmi et al. 2022a), with annotations provided by professional annotators. We also conduct thorough analysis on various characteristics of the translation (§5).

3 Experiments

In this section, we present various experiments. In §3.1, we describe several prompt selection strategies. In §3.2, we evaluate various GPT models in a zero-shot setup. In §3.3, we show results for high-resource language pairs, followed by results for low-resource and non-English pairs in §3.4. §3.5 provides document translation results. §3.6 examines the robustness of GPT models under domain shift. Finally, in §3.7, we discuss the potential of combining the benefits of GPT and NMT models.

3.1 Prompt Selection Strategies

It has been shown that the performance of LLMs can be enhanced through in-context learning by providing few labelled examples (prompts) in addition to the test input (Brown et al., 2020). This few-shot paradigm has demonstrated strong performance across multiple natural language processing (NLP) tasks (Ouyang et al., 2022; Goyal et al., 2022; Wei et al., 2022; Chowdhery et al., 2022). There has also been a series of recent works on in-context learning for machine translation (MT) with rather mixed results and various ways of shot selection. Recently, Zhang et al. (2023) use GLM-130B and show consistent but rather low correlation between the performance of MT and shot selection as compared to random. They use different features that show varying level of correlation to the performance. In the same spirit, Vilar et al. (2022) use different prompt selection schemes with PaLM-540B model with the main conclusion that shot selection is not necessarily better than random but they point out the importance of using high quality shots. In the same vein, Agrawal et al. (2022) use a much smaller model XGLM-7.5B and multiple selection criteria for the shots. They show that a combination of a retrieval and task metrics are consistently better than the random baseline across different translation directions.

In this paper, we explore prompt selection strategies along two dimensions: quality and relevance. Our pool to select few-shot examples is the cleaned WMT training data for each direction. This is obtained by filtering the full training data using language identification and length ratio. The size of the full and cleaned training data is shown in Table 12 for each direction. We do not use the development data from WMT shared tasks, despite its high quality, for shot selection to avoid any small chance of leaking information about the test set. In all cases, we test the performance with 0, 1 and 5 shots. In our preliminary experiments, we found that increasing beyond 5 shots did not result in any meaningful improvement. We show below how we select shots based on quality and relevance.

- *Quality*: To ensure high quality shots we sort our training data using LaBSE (Feng et al., 2020). We consider high quality shots that are randomly chosen from the top 1 million pairs as opposed to the full data. We also found it useful to select sentences that are longer than 50 tokens.

- *Relevance*: We consider relevant shots that are close to the input sentence. Based on preliminary experiments¹² we use the cosine distance between LaBSE embeddings as a measure of closeness. We always select relevant pairs from high quality ones (the top 1M pairs from LaBSE-scored training data). For computational efficiency, we adopt a two stage approach. We first use the input text and apply elastic search¹³ to retrieve the top 64 pairs then return the top 1 or 5 shots based on the LaBSE distance.

In the results, we refer to full random as *RR* (Random) while high quality are referred to as *QR* (Quality Random). The high quality shots selected through relevance are referred to as *QS* (Quality Selected).

3.2 Zero-Shot Translation Capabilities of GPT Models

We compare the general zero-shot translation capabilities of the three GPT models on four language pairs, in eight distinct translation directions. The selected languages were chosen with a focus on balancing representation. The languages include 1) German, which is one of the most represented non-English languages in GPT training data, 2) Russian, a large-scale non-Latin language, 3) Chinese, which represents a large-scale language with a script distinct from that of the majority of training data languages, and 4) French-German pair as non English-centric use case.

In this experiment, we compare the performance of three GPT models *text-davinci-002*, *text-davinci-003*, and *ChatGPT* with the top ranked systems in WMT22, as shown in Table 2. Remarkably, *text-davinci-002* shows lower performance, underperforming across all language pairs compared to the other two GPT models. On the other hand, *text-davinci-003* clearly shows better translation performance across all languages in this evaluation. Its zero-shot performance is comparable to the best performing WMT DE-EN system and outperforms the best ZH-EN WMT system.

ChatGPT shows a strong performance in the DE-EN language pair, while performing similarly to *text-davinci-003* when translating to English as well as French-German pairs. In terms of translat-

ing from English to other languages, *text-davinci-003* shows better performance than the other two GPT models. It is noteworthy that the translation between French and German, which is not English-centric, exhibits surprising competitiveness with state-of-the-art systems, despite the fact that the majority of the training data used for the GPT model is English-centric.

While both COMETkiwi and COMET-22 show relevant results, both lexical metrics (BLEU and ChrF) show consistent degradation with GPT models. This is consistent with similar findings from (Vilar et al., 2022) on PALM-540B model. We conducted human evaluation and more thorough analysis to further understand such results.

From these results, we can see that the three variants of GPT models exhibit different characteristics. However, the nature and extent of these differences remain unclear and require further investigation, depending on the availability of more information regarding the models, their training data, and their training methods. This superior performance of *text-davinci-003*, which is achieved in a zero-shot setting, motivates further investigation into the effect of few-shot in-context learning and shot selection strategies. We investigate these questions further in the sections below.

3.3 GPT Performance on High-resource Languages

Given the results in the previous section, we focus on evaluating *text-davinci-003* model, expanding the scope of the study to 18 language pairs and comparing its performance with that of a commercial system (Microsoft Translator) in addition to WMT SoTA systems. For consistency, in all the subsequent results, we use the term “GPT” to denote the *text-davinci-003* model, unless explicitly stated otherwise.

We experiment with various shot selection strategies: zero-shot, random (RR), quality (QR) and relevance selected (QS) prompts as described in §3.1. We report results for 1 and 5 shots along with the best WMT systems and MS-Translator.

Table 3 shows the performance of GPT *text-davinci-003* with few-shot setups on high-resource languages from WMT Testsets. With both reference and reference-less COMET scores, the model achieved impressive zero-shot results for all languages when translating into English. However, the few-shot configurations did not yield signifi-

¹²We also tried using COMET and embeddings based on a trained NMT model.

¹³Based on BM25 text retrieval.

System	COMET-22	COMETkiwi	ChrF	BLEU	COMET-22	COMETkiwi	ChrF	BLEU
	DE-EN				EN-DE			
WMT-Best	85.0	81.4	58.5	33.4	87.2	83.6	64.6	38.4
text-davinci-002	73.2	73.1	46.1	23.3	82.0	79.0	56.0	28.6
text-davinci-003	84.8*	81.2*	56.8	30.9	85.6*	82.8*	60.2*	31.8*
ChatGPT	84.8*	81.1	58.3*	33.4*	84.2	81.0	59.6	30.9
	ZH-EN				EN-ZH			
WMT-Best	81.0	77.7	61.1	33.5	86.7	82.0	41.1	44.8
text-davinci-002	74.1	73.1	49.6	20.6	84.0	79.0	32.1	36.4
text-davinci-003	81.6*	78.9*	56.0*	25.0	85.8*	81.3*	34.6	38.3
ChatGPT	81.2	78.3	56.0	25.9*	84.4	78.7	36.0*	40.3*
	RU-EN				EN-RU			
WMT-Best	86.0	81.7	68.9	45.1	89.5	84.4	58.3	32.4
text-davinci-002	77.5	76	58.7	34.9	85.4	80.9	51.6	25.1
text-davinci-003	84.8*	81.1*	64.6	38.5	86.7*	82.2*	54.0*	27.5*
ChatGPT	84.8*	81.0	66.5*	41.0*	77.6	70.4	41.1	19.0
	FR-DE				DE-FR			
WMT-Best	89.5	80.7	81.2	64.8	85.7	79.5	74.6	58.4
text-davinci-002	66.6	67.9	45.8	25.9	64.2	67.6	44.6	24.5
text-davinci-003	84.6	77.9	65.7*	42.5*	78.5	76.1	58.9	35.6
ChatGPT	84.7*	78.5*	65.2	42.0	81.6*	79.8*	60.7*	37.3*

Table 2: Zero-Shot evaluation results with three GPT models on 8 language pairs from WMT22 Testset. The best scores across different systems are marked bold. * denotes the best results among GPT systems.

cant improvements over the zero-shot setup. GPT surpassed both the WMT-Best and MS-Translator systems for DE-EN, JA-EN and ZH-EN language pairs, and almost matched the best systems for the other three language pairs. On the other hand, when translating from English to other languages, the few-shot setup consistently improved over the zero-shot setup, with most gains obtained from a single high-quality shot. GPT outperformed both the WMT-Best and MS-Translator systems for EN-JA and EN-ZH language pairs. We experimented with high quality shots with relevance scores (QS) for three languages (German, Russian and Chinese), but we observed no improvements over quality shots alone. This result emphasises the importance of fewer quality shots especially when translating from English. This difference in behavior is consistent with the observations that the critical role of demonstrations within in-context learning is to provide specifications of the output space (Min et al., 2022; Anonymous, 2023a), with a denser in-context learning signal being preferred when translating from English to other languages.

Similar to the Zero-Shot results from §3.2, we observe that the lexical metrics are showing consistent degradation with all GPT models and configurations.

3.4 GPT Performance on Low-resource and non English-centric Languages

We evaluated low-resource and non English-centric languages by conducting experiments with Icelandic and Hausa as two low-resource languages

and French and German as direct translation languages. Table 4 presents the results. The few-shot setup yields modest gains, especially when translating out of English. Similar to the high-resource case, most of the gains were obtained from a single high-quality shot. The systems for both low-resource languages did not surpass the WMT-Best systems. The DE-FR and FR-DE language pairs show remarkable results, as a single-shot setup outperforms the zero-shot setup significantly. This is consistent with the previous finding of translating from English to other languages; a more dense in-context signal is essential for direct translation as well, as it enables the model to generate in the correct language better than the zero-shot behavior. Both direct systems surpasses their commercial counterparts in terms of COMET scores, but they slightly trail behind the WMT-Best systems on the COMET-22 reference-based metric.

Similar to the high-resource language pairs, both lexical metrics (BLEU and ChrF) showed a significant and consistent degradation. To gain further insights into this, we conducted human evaluation and performed a more in-depth analysis as discussed in §4 and §5.

3.5 Document-Level MT with GPT

This section explores the application of GPT to document-level machine translation. Previous studies on MT with LLMs have mainly concentrated on sentence-level translation, with only a brief mention of document translation for transfer learning by Zhang et al. (2023). Document translation, on

System	COMET-22	COMETkiwi	ChrF	BLEU	COMET-22	COMETkiwi	ChrF	BLEU
	DE-EN				EN-DE			
WMT-Best	85.0	81.4	58.5	33.4	87.2	83.6	64.6	38.4
MS-Translator	84.7	81.0	58.5	33.5	86.8	83.4	64.2	37.3
GPT Zeroshot	84.8	81.2	56.8	30.9	85.6	82.8	60.2	31.8
GPT 1-Shot RR	84.9	81.3	56.1	30.4	86.1	83.0	60.7	31.9
GPT 1-Shot QR	84.9	81.3	56.7	31.1	85.8	82.8	60.7	32.4
GPT 5-Shot RR	85.2	81.5	56.5	31.2	86.5*	83.2 *	61.0	32.4
GPT 5-Shot QR	85.4*	81.5*	57.7	32.4	86.4	83.1	61.3*	33.2*
GPT 5-Shot QS	85.0	81.3	57.8*	32.5*	85.9	82.9	60.8	32.7
	CS-EN				EN-CS			
WMT-Best	89.0	82.5	79.3	64.2	91.9	85.3	68.2	45.8
MS-Translator	87.4	82.2	74.0	54.9	90.6	84.2	65.6	42.1
GPT Zeroshot	86.2	82.0	67.5	44.5	88.6	82.9	57.9	31.3
GPT 1-Shot RR	86.6	82.3	67.9	45.4	89.7*	84.0*	58.3	31.6
GPT 1-Shot QR	86.4	82.3	67.8	45.0	89.2	83.6	58.6	32.5
GPT 5-Shot RR	86.6	82.3	66.4	44.2	89.4	83.8	58.6	32.0
GPT 5-Shot QR	86.9*	82.5*	69.2*	47.5*	89.0	83.3	59.0*	32.9*
	JA-EN				EN-JA			
WMT-Best	81.6	80.3	49.8	24.8	89.3	85.8	36.8	27.6
MS-Translator	81.5	80.1	49.6	24.5	88.0	85.3	34.9	25.1
GPT Zeroshot	81.5	80.7	47.7	21.1	87.8	84.8	31.2	21.2
GPT 1-Shot RR	81.7	80.7	46.8	20.2	88.3	85.1	31.8	22.0
GPT 1-Shot QR	81.6	80.8	48.3*	22.1	88.4*	85.3	32.2*	22.5*
GPT 5-Shot RR	82.0*	80.9*	48.2	22.4*	88.2	85.4*	31.7	21.4
GPT 5-Shot QR	81.8	80.8	47.2	21.0	88.2	85.3	31.1	21.6
	ZH-EN				EN-ZH			
WMT-Best	81.0	77.7	61.1	33.5	86.7	82.0	41.1	44.8
MS-Translator	80.4	77.6	57.7	27.9	86.1	81.4	43.1	48.1
GPT Zeroshot	81.6*	78.9*	56.0*	25.0*	85.8	81.3	34.6	38.3
GPT 1-Shot RR	80.9	78.2	55.2	24.2	86.7	81.8	38.7	42.8
GPT 1-Shot QR	81.2	78.8	55.3	24.2	86.1	81.5	35.5	38.8
GPT 5-Shot RR	81.1	78.8	55.0	24.4	87.0	82.0	37.1	41.3
GPT 5-Shot QR	81.1	78.7	54.7	23.8	87.0*	82.2*	39.8*	43.7*
GPT 5-Shot QS	81.0	78.5	55.5	24.6	86.2	81.5	38.3	41.8
	RU-EN				EN-RU			
WMT-Best	86.0	81.7	68.9	45.1	89.5	84.4	58.3	32.4
MS-Translator	85.2	80.7	68.3	43.9	87.4	82.9	58.1	33.1
GPT Zeroshot	84.8	81.1	64.6	38.5	86.7	82.2	54.0	27.5
GPT 1-Shot RR	84.1	80.6	63.3	37.9	86.4	81.9	54.3	28.1
GPT 1-Shot QR	84.9	81.2*	65.4*	40.1	86.9	82.4	53.8	27.5
GPT 5-Shot RR	84.9	81.2*	63.9	39.0	86.8	82.3	54.3	27.9
GPT 5-Shot QR	84.9	81.0	65.4*	40.0	87.0*	82.4*	54.4*	28.2*
GPT 5-Shot QS	85.0*	81.2*	65.3	40.2*	86.4	82.2	54.4*	28.0
	UK-EN				EN-UK			
WMT-Best	86.0	81.5	67.3	44.6	88.8	83.4	59.3	32.5
MS-Translator	83.5	79.7	65.3	42.4	86.1	81.9	56.1	28.2
GPT Zeroshot	83.5	80.1	59.8	34.8	83.7	79.5	49.6	21.1
GPT 1-Shot RR	83.5	80.3	60.3	35.6	84.7	80.2	50.1	21.2
GPT 1-Shot QR	83.8	80.3	61.4	37.5	85.1	80.5	50.5	21.9
GPT 5-Shot RR	83.6	80.3	58.8	34.4	85.4*	80.8*	50.9*	22.6*
GPT 5-Shot QR	83.9*	80.3*	62.1*	38.4*	85.4	80.6	50.6	22.1

Table 3: Zero-Shot and Few-Shots evaluation results with GPT(text-davinci-003) on high resource languages from WMT Testsets. The best scores across different systems are marked bold. * denotes the best results among GPT systems.

System	COMET-22	COMETkiwi	ChrF	BLEU	COMET-22	COMETkiwi	ChrF	BLEU
	IS-EN				EN-IS			
WMT-Best	87.0	81.4	62.3	41.7	86.8	81.8	59.6	33.3
MS-Translator	85.9	80.3	62.8	40.5	84.3	80.2	56.8	28.7
GPT Zero-shot	82.1	78.7	55.6	31.9	76.3	74.0	43.5	15.9
GPT 1-Shot RR	84.1	80.2	57.8	34.7	77.0	74.6	43.7	15.3
GPT 1-Shot QR	83.5	79.7	56.7	33.3	77.4	75.1	44.5	16.2
GPT 5-Shot RR	84.4*	80.4*	58.1*	35.0*	77.9*	75.2*	45.1*	16.8*
GPT 5-Shot QR	84.2	80.2	58.0	35.2	76.0	74.1	44.1	16.3
	HA-EN				EN-HA			
WMT-Best	80.0	74.5	48.7	21.0	79.8	61.5	51.1	20.1
MS-Translator	73.3	68.5	43.4	16.2	72.5	57.2	38.4	10.3
GPT Zero-shot	76.1	73.1	45.5	17.3	73.3	58.6	38.4*	9.4*
GPT 1-Shot RR	75.7	72.7	45.7	17.3	74.0	59.0	38.4*	8.8
GPT 1-Shot QR	78.1	74.4	47.5*	19.1*	74.1*	59.7*	37.8	8.9
GPT 5-Shot RR	75.5	72.2	45.9	17.8	72.1	57.7	36.0	8.0
GPT 5-Shot QR	78.2*	74.5*	47.5*	18.9	72.6	58.5	36.9	8.5
	FR-DE				DE-FR			
WMT-Best	89.5	80.7	81.2	64.8	85.7	79.5	74.6	58.4
MS-Translator	85.4	78.9	67.5	45.3	82.7	79.0	65.0	42.0
GPT Zero-shot	84.6	77.9	65.7	42.5	78.5	76.1	58.9	35.6
GPT 1-Shot RR	86.1	79.6	65.1	41.0	83.1	80.5	60.3	36.9
GPT 1-Shot QR	86.4	80.0	67.0	43.9	83.2	80.8	61.2	38.1
GPT 5-Shot RR	86.6	80.0	65.2	41.6	83.6*	80.9*	60.1	37.1
GPT 5-Shot QR	86.7*	80.2*	67.7*	44.8*	83.2	80.7	62.1*	39.3*

Table 4: Zero-Shot and Few-Shots evaluation results with GPT (text-davinci-003) on low resources and non-English centric translation directions from WMT Testsets. The best scores across different systems are marked bold. * denotes the best results among GPT systems.

the other hand, has received considerable attention for transformer models, with inconclusive findings as to the efficacy of the additional context. Sun et al. (2020) challenge some of the prior studies and demonstrate that a simple modification of the training to vary the document lengths can significantly enhance the performance of a standard transformer architecture for document-to-document translation. We hypothesize that GPT can excel at document-to-document translation, as it is trained on large contexts. Moreover, translating entire documents can reduce the number of API calls and thus improve the computational efficiency and latency. We argue that document translation improvements may need better metrics to capture its potential. Hence, in this section, we report results using doc-BLEU (Liu et al., 2020) and doc-COMET metrics as described in §2.4, in addition to our sentence-level metrics.

Evaluating Document-level Translations A document-level translation does not necessarily keep sentence-level alignments intact. We try to prompt GPT to keep sentence-level alignment intact by emphasizing sentence separation in the prompt. Our prompt template can be found in Figure 18 from the appendix. However, we find that we still need to restore sentence-level alignment with source for some of the documents in the test set. In all cases, we find two types of mismatch

that we need to restore. First, sentences that were written in the source over two lines and their translation is one line. In that case, we insert a new line break to match the position of the new line break in the source sentences. Second, sentences that are skipped. In that case, we replace empty lines at the end of the documents with empty lines in place of the skipped sentences.

We need to restore sentence-level alignment to calculate metrics that were mainly developed for sentence-level evaluation such as the COMET-22 and COMETkiwi models that we are using. We also follow (Liu et al., 2020) and report SacreBLEU calculated on the document level. For neural network based metrics, we extend the COMETkiwi model for document level evaluation as described in §2.4.

Experiment 1: We conduct a series of experiments on zero-shot MT using News Commentary dataset, varying the window length from 1 (sentence-level) to 32 in powers of two¹⁴. Table 5 shows that increasing the window size leads to improvements across all metrics. However, the gains in lexical metrics (BLEU and ChrF) are larger than in neural metrics (COMET-22 and COMETkiwi). The document-based metric (doc-BLEU and doc-COMETkiwi) also exhibit similar improvements

¹⁴The context is limited to the document if the window length exceeds the document length.

to the sentence-based metric. Remarkably, as the window size grows, the performance surpasses MS-Translator model and approaches the WMT-Best systems. This is consistent with the findings of Sun et al. (2020) for conventional MT models. The table also displays the total number of requests for each window size. We observe that the number of requests decreases dramatically as the window size increases, while the performance either improves significantly or remains relatively stable depending on the metric used. Therefore, this document-level setup achieves high efficiency without compromising the quality.

Experiment 2: The second set of experiments investigates few-shot translation for the document setting. Following the sentence experiments, we focus on using 5-shots. We also use the News Commentary dataset, which has document-level annotations. Table 6 summarizes the results. The first two rows show the best WMT22 and the MS-Translator results for reference.

The following rows are named GPT-XX-YY where XX stands for the scope of translation (sentence or document) and YY stands for the source of the 5 shots (QR, DR, DF or DH as explained below). Rows GPT-Sent-QR and GPT-Sent-DR show results for sentence-level translation. The former uses the same quality-based shots from Table 3, while the latter uses 5 randomly selected shots from the document set, excluding the test data. We perform document translation (referred to as Doc in the table) for a window of 10 sentences in the following rows. Rows named GPT-Doc-QR and GPT-Doc-DR use the same shots as the sentence case. For row GPT-Doc-DF, we select a random document from the document data pool and use the first 5 sentences of the document as shots (i.e. document first DF). For row GPT-Doc-DH, we store GPT outputs in history for use in shots. We translate the first document 0-shot, and the subsequent documents 5-shots. For shot selection, we pick a random document from the previously translated documents and use the first 5 sentences as shots (i.e. document history DH). The results indicate that document translation outperforms sentence translation across metrics. However, while few-shots yield some consistent gains for sentence translation, this is not the case for document translation. This may be explained by the fact that document translation provides enough context, making few-shots redundant. It can be also observed from the table

that the Doc-COMETKiwi shows more gain than the sentence level metrics but this might need more indepth analysis to verify.

3.6 Robustness Toward Domain Shift

We use the WMT datasets to examine how domain shift affects the performance of GPT models on German and Chinese, both from and to English. The WMT22 datasets span four domains: Conversational, News, e-Commerce and Social. Table 7 shows the scores of the four domains on the WMT testsets.

GPT achieves remarkable improvements on the conversational domain for DE-EN, ZH-EN and EN-ZH, as evidenced by both COMET and lexical scores (BLEU and ChrF). This contrasts with previous observations where lexical scores consistently deteriorate with GPT models.

GPT performs comparably to the other systems on the news domain according to COMET scores for all directions. It surpasses both other systems on DE-EN, while slightly trailing behind on EN-DE. For ZH-EN and EN-ZH, GPT exceeds MS-Translator, but falls slightly short of WMT-Best system. However, GPT scores significantly lower in terms of BLEU metric for both ZH-EN and EN-DE.

GPT clearly outperforms both systems on ZH-EN and matches WMT-Best on DE-EN for the e-commerce domain. It slightly lags behind on other directions. In this domain, we observe consistent lower scores in BLEU metric for all directions even for ZH-EN which outperforms significantly on both COMET metrics.

GPT outperforms both systems on DE-EN for the social domain. However, on ZH-EN and EN-ZH, GPT only surpasses them on COMETkiwi, while showing lower BLEU score for all directions with a significant difference in ZH-EN which exhibits substantial gains on COMETkiwi.

The results demonstrate GPT’s robust translation capabilities across different domains and languages. It performs well on DE-EN, ZH-EN and EN-ZH for all domains. However, we observe a discrepancy on lexical metrics for ZH-EN and DE-EN on News and Social domains even with GPT’s high performance on those languages. We conduct a further examination of the ZH-EN results to gain more insights. We find that the news from Chinese outlets follows a more templatic style especially in the prefix of the news. For NMT systems that are trained

System	COMET-22	COMETkiwi	Doc-COMETkiwi	ChrF	BLEU	Doc-BLEU	GPT Requests
DE-EN							
WMT-Best	85.0	81.4	79.9	58.5	33.4	35.2	–
MS-Translator	84.7	81.0	79.5	58.5	33.5	35.2	–
GPT Sent ZS	84.8	81.2	79.5	56.8	30.9	32.3	1984
GPT Doc ZS w=2	85.1	81.4*	80.0	57.8	32.6	34.4	1055
GPT Doc ZS w=4	85.2*	81.3	80.2*	57.9	32.8	34.5	607
GPT Doc ZS w=8	85.1	81.2	80.2	57.9	33.0	34.7	401
GPT Doc ZS w=16	85.2	81.2	80.2	58.0*	33.1*	34.8*	310
GPT Doc ZS w=32	85.1	81.2	80.2	57.9	33.1	34.8	274
EN-DE							
WMT-Best	87.2	83.6	83.1	64.6	38.4	40	–
MS-Translator	86.8	83.4	83	64.2	37.3	38.8	–
GPT Sent ZS	85.6	82.8	82.2	60.2	31.8	33.1	2037
GPT Doc ZS w=2	86.1	82.7	82.4	60.9	32.8	34.4	1058
GPT Doc ZS w=4	86.3	82.6	82.6	61.3	33.6	35.2	579
GPT Doc ZS w=8	86.4	82.6	82.6	60.9	33.4	35.2	349
GPT Doc ZS w=16	86.5*	82.6*	82.6*	61.3*	34.2*	36.1*	235
GPT Doc ZS w=32	86.4	82.6	82.7	61.3	34.1	36.1	187

Table 5: Evaluation results of document-level translation with GPT on DE<>DE WMT22 testset. The table shows the effect of increasing context length w in document to document translation with a zero-shot setting.

System	COMET-22	COMETkiwi	Doc-COMETkiwi	ChrF	BLEU	Doc-BLEU
DE-EN						
WMT-Best	85.0	81.4	79.9	58.5	33.4	35.2
MS-Translator	84.7	81.0	79.5	58.5	33.5	35.2
GPT-Sent-QR	85.4*	81.5*	80.2	57.7	32.4	34.0
GPT-Sent-DR	84.9	81.4	80.0	55.3	29.6	31.3
GPT-Doc-QR	85.2	81.2	80.2	58.1	33.3	35.0
GPT-Doc-DR	85.2	81.3	80.5*	57.6	32.7	34.3
GPT-Doc-DF	85.1	81.3	80.4	57.3	32.6	34.1
GPT-Doc-DH	85.3	81.2	80.3	58.2*	33.5*	35.1*
EN-DE						
WMT-Best	87.2	83.6	83.1	64.6	38.4	40.0
MS-Translator	86.8	83.4	83.0	64.2	37.3	38.2
GPT-Sent-QR	86.4	83.1	82.7	61.3	33.2	34.8
GPT-Sent-DR	86.7	83.5*	83.0*	61.2	32.8	34.3
GPT-Doc-QR	86.6	83.0	82.7	61.6	34.0	35.7
GPT-Doc-DR	86.9	83.0	82.9	61.9	34.4	36.1
GPT-Doc-DF	87.0*	83.1	82.9	62.0*	34.5*	36.2*
GPT-Doc-DH	86.6	82.9	82.8	61.7	33.9	35.7

Table 6: Effect of shot selection for document-level translation on WMT22 DE<>EN testset.

System	COMET-22	COMETkiwi	ChrF	BLEU	COMET-22	COMETkiwi	ChrF	BLEU
<i>Combined All Domains</i>								
		DE-EN				EN-DE		
WMT-Best	85.0	81.4	58.5	33.4	87.2	83.6	64.6	38.4
MS-Translator	84.7	81.0	58.5	33.5	86.8	83.4	64.2	37.3
GPT	85.4	81.5	57.7	32.4	86.4	83.1	61.3	33.2
		ZH-EN				EN-ZH		
WMT-Best	81.0	77.7	61.1	33.5	86.7	82.0	41.1	44.8
MS-Translator	80.4	77.6	57.7	27.9	86.1	81.4	43.1	48.1
GPT	81.1	78.7	54.7	23.8	87.0	82.2	39.8	43.7
<i>Conversational Domain</i>								
		DE-EN				EN-DE		
WMT-Best	85.7	81.4	54.9	35.1	89.1	83.3	67.7	42.8
MS-Translator	85.1	81.0	55.2	35.3	88.8	83.1	67.3	40.7
GPT	86.1	81.5	55.0	35.6	88.5	83.4	62.9	35.7
		ZH-EN				EN-ZH		
WMT-Best	81.6	77.7	48.8	25.6	86.9	81.3	36.7	37.6
MS-Translator	80.6	77.4	46.7	22.9	87.6	81.3	42	44.0
GPT	82.0	78.1	48.0	26.3	88.9	82.4	40.5	41.9
<i>News Domain</i>								
		DE-EN				EN-DE		
WMT-Best	84.9	82.0	58.8	31.4	87.0	84.5	65.6	37.8
MS-Translator	84.7	81.9	59.0	31.6	87.0	84.3	65.3	36.8
GPT	85.3	82.3	58.7	31.4	85.9	83.7	62.4	31.7
		ZH-EN				EN-ZH		
WMT-Best	82.1	79.6	62.2	31.3	87.6	83.1	45.6	51.7
MS-Translator	81.8	80.0	59.7	28.2	87	82.3	48.2	53.8
GPT	81.7	80.2	56.9	23.3	87.2	82.5	42.2	48.7
<i>e-Commerce Domain</i>								
		DE-EN				EN-DE		
WMT-Best	85.6	81.2	61.3	35.3	88.7	84.1	65.3	38.4
MS-Translator	85.3	80.9	61.3	35.2	88.2	83.8	64.9	38.1
GPT	85.6	81.2	60.5	34	87.6	83.6	62	34
		ZH-EN				EN-ZH		
WMT-Best	77.6	75.1	52.2	22.2	88.0	83.0	40.8	43.5
MS-Translator	77.8	75.0	51.5	20.3	87.8	82.6	41.9	46.6
GPT	79.0	76.7	49.4	18.7	87.9	82.8	40.3	42.9
<i>Social Domain</i>								
		DE-EN				EN-DE		
WMT-Best	84.1	80.9	56.5	32.6	84.0	82.3	59.8	35.8
MS-Translator	83.7	80.5	56.2	32.5	83.2	82.2	59.1	34.6
GPT	84.5	81.1	54.7	29.9	83.6	81.7	57.7	32.5
		ZH-EN				EN-ZH		
WMT-Best	83.0	78.2	69.4	46.8	84.2	80.7	36.4	40.2
MS-Translator	81.4	78.0	62.3	34.6	82.0	79.4	36.1	40.9
GPT	81.9	79.7	57.4	28	84.0	81.0	34.2	37.7

Table 7: Evaluation results of DE \leftrightarrow EN and ZH \leftrightarrow EN translations across four domains.

heavily on similar data, it is easier to reproduce the same patterns, e.g., WMT-Best scores 31.3 BLEU. For more general commercial scale systems that are trained on much larger and diverse data, it is harder to produce the same exact patterns, as such, MS-Translator scores 28.2 BLEU. For GPT, which is mostly trained on English, it is harder to get the lexical matches and it mostly produces the English news style, scoring 23.3 BLEU. However, COMET-22, using the same reference, seems to provide a more robust signal with all three systems being almost on the same level. This confirms the robustness of neural metrics across domains (Freitag et al., 2022) and GPT’s ability to handle diverse domains while being more robust towards *parallel data biases*, which we will explore further in §5. Therein, we show that in general, GPT performs better in cases where the input *resonates* with the noisy parts of parallel data.

3.7 Hybrid GPT and NMT Translation

To explore the possibility of leveraging the strong performance of GPT on various languages, we propose and evaluate several hybrid approaches that combine the strengths of NMT and GPT systems. The basic idea is to use Microsoft Translator (MS-Translator) system as the primary translation system, and then use GPT as a fallback system when the quality of MS-Translator is unsatisfactory.

We use COMETkiwi as the quality estimation model and COMET-22 as the performance evaluation metric. We first establish an upper bound by selecting the best translation from either systems according to COMETkiwi, which we call the “Max-Routing” approach. Then we experiment with a more practical approach where we use GPT only when the COMETkiwi score of MS-Translator falls below a predefined threshold. In this experiment, we set the threshold to the 50-th percentile of the COMETkiwi scores of MS-Translator, meaning that we use GPT for any translation that has a COMETkiwi score lower than the median MS-Translator COMETkiwi score, which can be easily estimated from previous translation requests.

Figure 1 presents the results of our experiments on 12 language pairs. Firstly, we observe that in all language pairs, the “Hybrid Max-Routing” approach consistently achieves the highest COMET-22 scores, surpassing both the individual systems. “Hybrid Max-Routing” achieves a maximum gain of 1.6 Comet-22 points in the EN-UK language pair

which is not among top performance language for GPT. This indicates that combining the strengths of NMT and GPT systems can lead to a significant improvement in translation quality.

Next, we compare the performance of the individual systems. In general, MS-Translator achieves higher scores than GPT on most language pairs, which is expected given that MS-Translator is an NMT system specifically optimized for translation tasks. However, GPT outperforms MS-Translator on certain language pairs, such as DE-EN, EN-JA, and EN-ZH. This suggests that GPT can be a valuable fallback system in cases where the quality of the primary system is unsatisfactory.

We also compare the performance of the two hybrid approaches. The “Hybrid Max-Routing” approach achieves slightly higher scores than the “Hybrid Threshold” approach on most language pairs, indicating that routing to GPT only when the quality of MS-Translator falls below a certain threshold may not always be the optimal strategy. However, the “Hybrid Threshold” approach still achieves comparable results to the upper bound on all language pairs, while using GPT for only 50% of the instances. This suggests that it can be a more practical approach in scenarios where computational resources are limited.

Figure 2 demonstrates that hybrid approaches achieve larger and more consistent improvements than shots selection across all languages and directions. Figure 3 zooms in on the high-performing DE-EN and EN-DE systems. The hybrid system outperforms both WMT-Best and MS-Translator systems in both directions, even though the GPT system only outperforms them in DE-EN with 5-shot setup.

In summary, our experiments demonstrate the potential of combining NMT and GPT systems to improve machine translation quality. The results suggest that a hybrid approach that uses GPT as a fallback system can achieve higher performance than either individual systems. Future research can explore more advanced techniques that can leverage the strengths of both systems and optimize the hybrid approach.

4 Human Evaluation and Analysis

We use source-based sentence-level contrastive Direct Assessment + Scalar Quality Metric (contrastive DA+SQM; Akhbardeh et al. 2021, Kocmi et al. 2022a) to perform human evaluation of the




Figure 1: Comparing COMET-22 scores of hybrid MS-Translator and GPT systems with GPT and MS-Translator systems.

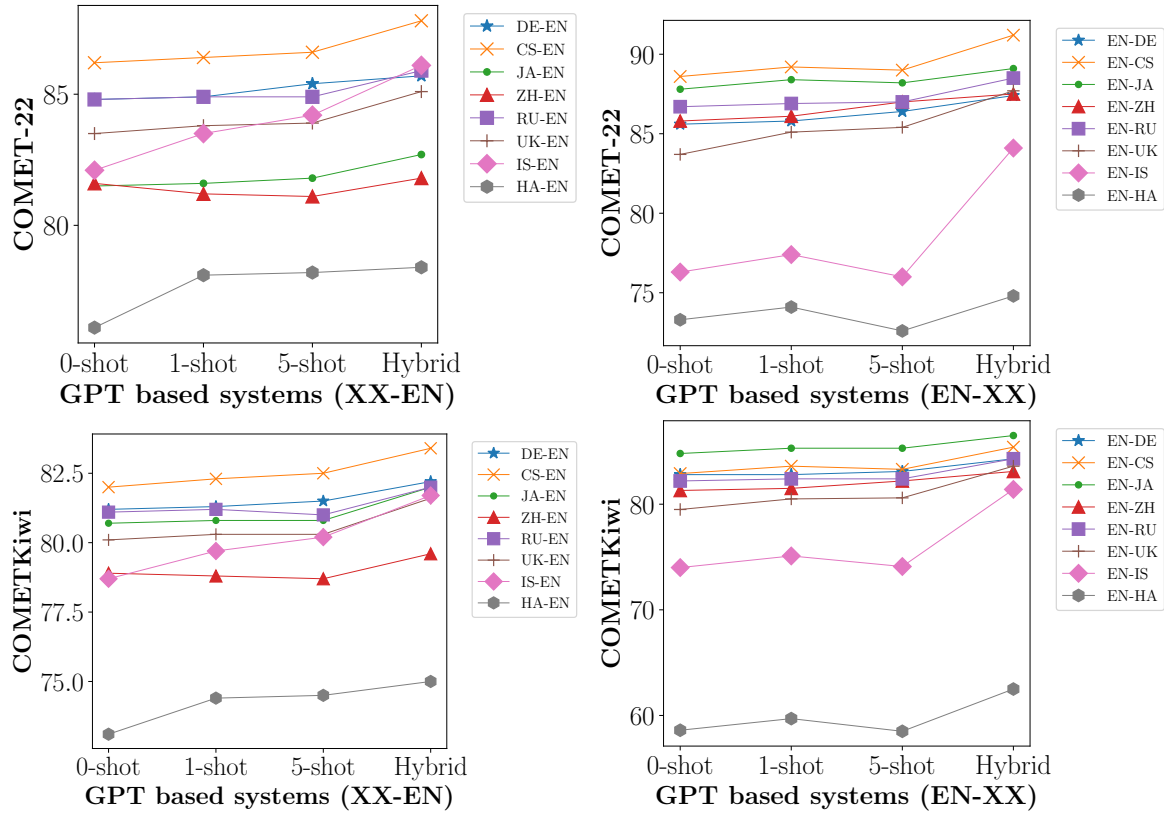


Figure 2: COMET-22 and COMETKiwi scores for GPT based systems with different approaches.

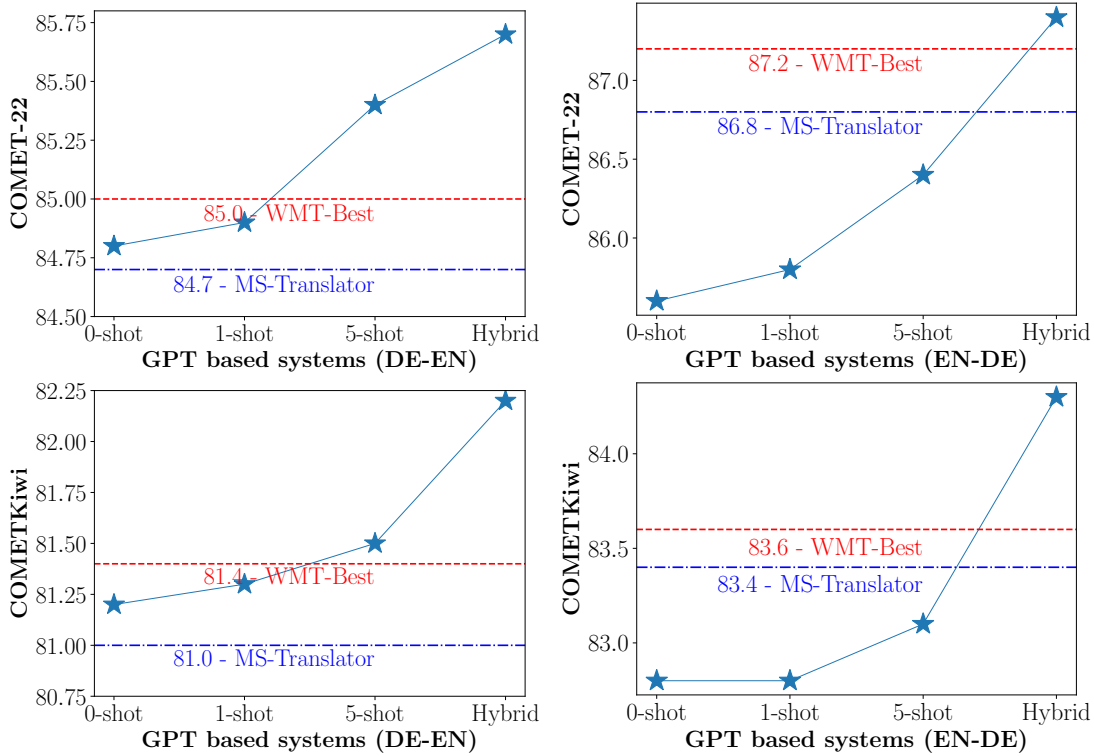


Figure 3: COMET-22 and COMETKiwi scores for GPT based systems compared to WMT-Best and MS-Translator systems to translate between English (EN) and German (DE).

WMT-Best systems from Table 1 and GPT with 5 shots QR as shown in Table 3. For each language pair, we randomly sample 425 non-identical translation item pairs and have them annotated with the contrastive DA+SQM annotation method by 5 distinct professional translation experts per language pair. Figure 4 and Figure 5 report aggregated human and corresponding COMETKiwi scores. Surprisingly, GPT outperforms Best-WMT systems on CS-EN, ZH-EN, EN-ZH and DE-FR, and achieves comparable results on most of the high-resource languages. On the other hand, the two low-resource languages, Hausa and Icelandic, lag behind significantly. Full details of the scores can be found in the appendix, Table 13.

We observe that the human evaluation results are highly consistent with the COMETKiwi results. This highlights the importance of neural reference-less metrics for evaluating MT in general and this family of models in particular. As we have seen in the previous results, all lexical metrics fail to capture the strong performance of GPT and exhibit lexical and reference bias. While we believe quality estimation is becoming more essential for MT in general, it is reassuring to know that COMETKiwi performs well on GPT models as well as on NMT models. Moreover, as shown on

Figure 6, highly performing GPT languages pairs demonstrate higher win rate which is reflected on both Human Evaluation results and COMETKiwi scores.

We conducted a manual analysis of human-evaluated GPT translations for English-Japanese and Japanese-English directions to identify their strengths and weaknesses. In Table 14 in the appendix, we present some of the observed characteristics along with examples of GPT and WMT outputs. A notable characteristic is that GPT performs better and more robustly than WMT for source sentences that are erroneous, short, or colloquial. We found that GPT can handle misspellings or unclosed quotation marks and produce translations that do not omit any semantic information. Moreover, GPT can generate reasonable translations for partial or incomplete colloquial source sentences while WMT-Best often adds or omits content. However, GPT tends to produce unnatural translations for sentences with uncommon or complex expressions.

5 GPT Translation Characteristics

In this section, we try to comprehensively analyze the characteristics of the translations obtained from GPT. Our goal here is to better differentiate GPT




Figure 4: Human Evaluation of WMT-Best Systems vs GPT translations for non-English to English




Figure 5: Human Evaluation of WMT Best Systems vs GPT translations to non-English

translations from its NMT counterpart.

5.1 Situating GPT Translations

We posit that there are two key biases due to which the computation of translation done by LLMs might be different from the same computation done by NMT models, namely the *Parallel Data Bias* and the *Language Modeling Bias*.

Parallel Data Bias: Compared to NMT models trained on parallel data, which is typically web-mined (and noisy), LLMs such as GPT are trained on monolingual data only with no explicit supervisory signal for the translation task. This creates interesting implications on the nature of the

emergent computational abilities leveraged for our task of interest, translation. First, not using parallel data might imply that LLMs are protected against the noise associated with parallel data, which leads to problems such as the memorization of noisy/atypical (Raunak et al., 2021) or low-quality samples (Raunak and Menezes, 2022) and biases towards particular language characteristics predominant in the parallel data (Garcia et al., 2023). These *parallel data biases* can also manifest in the form of long-tailed errors such as the translations of physical units or currencies (Raunak et al., 2022), owing to a preponderance of such incorrect token pairings in the parallel data. On the other hand, the lack of explicit supervisory signals




Figure 6: Human Evaluation: GPT Win Rates (%) based on Item Scores per language pair.




Figure 7: **Fluency Comparisons** for the X-E language pairs. On 7 out of 8 language pairs, GPT translations obtain lower perplexity, thereby producing more fluent translations. The magnitude of the difference is higher for Zh-En and Ja-En language pairs.




Figure 8: **Comparisons of Punctuation Insertions** for the X-E language pairs. On 8 out of 8 language pairs, GPT translations show a greater bias towards inserting unsupported end of sentence markers in translation.

for the task could also mean that LLM based translations might not track the desired characteristics of translations such as faithfulness to the source as well as the NMT models trained with explicit teacher-forced supervision (Anonymous, 2023b).

Language Modeling Bias: Despite the impressive performance of in-context learning, constraining LLM behavior to explicitly follow the specifications of a desired task is a non-trivial problem. Analyses of in-context learning have revealed how the *implicit* zero-shot performance of LLMs might be higher than their *observed* zero-shot performance, with the demonstrations within in-context learning themselves providing only limited learning signals (Min et al., 2022; Kojima et al., 2022; Anonymous, 2023a). A direct implication of these

results for translation is that the demonstrations used for in-context learning might *fail* to override the underlying computational bias of language modeling which is likely to favor greater fluency at the cost of adequacy. Such *language modeling bias* might also introduce undesirable artifacts, e.g., punctuation insertions, acronym expansions, world knowledge insertion, etc. in the translations which could cause it to veer off from a faithful cross-lingual representation of the input.

In the next subsection, we propose properties along which finer-grained characteristics of GPT translations could be enumerated. These measures are designed to provide indirect measurements of the language modeling bias as well as the parallel data bias, which could allow a better differentiation of GPT translations against translations from

Sequence Type	Translation Instance	Phenomenon
Source	Bis auf die E 95 02 wurden alle Lokomotiven zerlegt .	Non-Monotonicity (NM)
MS Translator	With the exception of E 95 02, all locomotives were dismantled .	
GPT	All locomotives were dismantled except for the E 95 02.	
Source	Order ist sie ganz aus dem Sortiment genommen?	Fluency (F)
MS Translator	Or is it completely removed from the range?	
GPT	Or has it been completely removed from the range?	
Source	Sehen Sie bitte im Screenshot was der Kollege geschrieben hat	Punctuation Insertion (PI)
MS Translator	Please see in the screenshot what the colleague wrote	
GPT	Please see the screenshot for what the colleague wrote.	
Source	Die Email zur Stornierung wurde am 26.12.#NUMBER# versendet.	Dropped Content (USW)
MS Translator	The cancellation email was sent on 26.12.#NUMBER#.	
GPT	The cancellation email was sent on December 26th.	
Source	"We won't accept the CAA and that is for sure.	Inserted Content (UTW)
MS Translator	"我们不会接受 CAA ，这是肯定的。	
GPT	"我们不会接受《公民法》，这是肯定的。	

Table 8: **Illustrated Examples of the Phenomena** as described in Section 5. The origin of these differences between translations lie in the computational mechanism leveraged for translations: When controlled for quality, higher translation non-monotonicity suggests a more abstractive computation used for obtaining the translations. Similarly, Fluency, Punctuation Insertion, Dropped and Inserted Content measure different translation characteristics.




Figure 9: **Comparisons of Unaligned Source Words** for the X-E language pairs. GPT Translations consistently incur greater number of unaligned source words.




Figure 10: **Comparisons of Unaligned Translation Words** for the X-E language pairs. GPT Translations consistently incur greater number of unaligned target words.

NMT systems. We first discuss the measurements designed to elicit artifacts associated with the language modeling bias.

5.2 Language Modeling Bias Artifacts

We propose and use five measurements over the test sets to quantitatively explore language modeling bias, in order to enumerate the differences in translations obtained from traditional NMT systems and GPT. Below, we describe the properties as well as the algorithms used for quantifying them (corresponding illustrative examples of the phenomena are presented in Table 8):

1. **Translation Non-Monotonicity (NM):** We aim to measure how closely the translation

tracks the source sentence. A more paraphrastic or a less literal translation is likely to deviate from a close tracking of the source word order (across language pairs). We use the non-monotonicity metric proposed in [Schioppa et al. \(2021\)](#), which computes the deviation from the diagonal in the word to word alignment as the non-monotonicity measure. This measurement could also be interpreted as a normalized measure of alignment crossings, which has been shown to correlate with translation non-literality ([Schaeffer and Carl, 2014](#)). This measurement has also been used in [Anonymous \(2023b\)](#) for investigating translation literalness.




Figure 11: **Comparisons of Translation Non-Monotonicity** for the X-E language pairs. GPT Translations consistently score higher on the non-monotonicity of translations.




Figure 12: **Comparisons of Punctuation Insertions** for the E-X language pairs. On 8 out of 8 language pairs, GPT translations obtain higher scores.

2. **Translation Fluency (TF)**: We measure translation fluency using a strong, independently trained language model ('gpt2-large', Radford et al. (2019)). We restrict this measurement to X-E direction, since GPT-2 has only been trained on English text (Radford et al., 2019).
3. **Punctuation Insertion (PI)**: Language modeling bias can prefer one mode of sentence completion in contrast to others. This can reveal itself in the presence of not well-formed inputs such as sentences that do not end with typical end of sentence markers (comma, period and exclamation). We measure the fraction of input sentences for which the translation contains an end of sentence marker but the source does not. The insertion of an end of sentence marker in such instances is inadequate for translation, a task which strives for




Figure 13: **Comparisons of Unaligned Source Words** for the E-X language pairs. GPT Translations, on average, incur greater number of unaligned source words.




Figure 14: **Comparisons of Unaligned Translation Words** for the E-X language pairs. GPT Translations consistently incur greater number of unaligned target words.

bitext equivalency.

4. **Unaligned Source Words (USW)**: We measure the number of source words left unaligned in a word to word alignment obtained over the source and output translations. When controlled for quality, a more paraphrastic translation is likely to contain more words that do not align with the words in the source sentence. This measurement was used in Anonymous (2023b) as a measure of translation literalness and we use it similarly to obtain a measurement of content that is dropped in a translation – an untranslated word or phrase in a source sentence is likely to find no alignments in the output. For obtaining word to word alignments, we use a multilingual-bert based aligner (Devlin et al., 2019; Dou and Neubig, 2021).
5. **Unaligned Translation Words (UTW)**: We




Figure 15: **Comparisons of Translation Non-Monotonicity** for the E-X language pairs. GPT Translations score higher on translation non-monotonicity for 3 out of 4 high-resource language pairs.

measure the number of unaligned words in the translation using the same word to word alignments as in the previous measurement. This indicates the presence of words that have no support in the source and is included to measure words that are potentially inserted in the translation without any basis in the input.

We collect the measurements on these properties over the test sets for all the language pairs under investigation. We compare MS Translator with GPT throughout. We report the results in the next section and present our analysis grouped by the translation direction.

5.3 X-E Translation characteristics

Figures 7, 8, 9, 10 and 11 represent the comparisons of GPT translations against MS Translator for X-E language pairs. Figure 7 shows that GPT translations obtain lower perplexities, thereby demonstrating greater fluency. Figure 8 shows that GPT translations suffer from the problem of punctuation insertion with a much higher frequency than MS translator. We attribute this to the language modeling bias, which would prefer to generate a well-formed sentence, even if such well-formedness is unsupported in the input. Figure 9 shows that GPT translations incur slightly higher number of unaligned source words on 7 out of 8 eight language pairs. Greater unaligned source words would imply either the presence of greater paraphrasticity in the translations or greater inadequacy (dropped or inserted content). Figure 10 shows that GPT translations incur almost similar number of unaligned target words, suggesting that the GPT

Lang-Pair	System	PI ↓	NM ↓	USW ↓	UTW ↓
De-Fr	MS-Translator	3.61	17.68	10.91	25.39
	GPT	42.98	17.21	11.42	25.11
Fr-De	MS-Translator	2.63	14.63	21.87	14.63
	GPT	60.30	14.52	21.77	14.52

Table 9: Translation Characteristics Comparisons for **De-Fr and Fr-De**: GPT shows a much higher tendency to add an end of sentence marker into the translation when it is absent in the source.

translations are similarly adequate in terms of potential insertions. Another measurement, presented in Figure 11 shows that GPT translations are more non-monotonic than its NMT counterpart.

5.4 E-X Translation characteristics

Figures 12, 13, 14 and 15 represent the comparisons of GPT translations against MS Translator for E-X language pairs. Figure 12 shows that similar to X-E translations, GPT E-X translations also suffer from a higher frequency of punctuation insertions. The magnitude of the difference however, is smaller than the X-E translations, suggesting a weaker language modeling bias for these languages. Figure 13 shows that in general, GPT translations incur greater number of unaligned source words than its NMT counterpart. Figure 14 shows that the number of unaligned translation words for GPT translations do not differ greatly from MS Translator. Similarly, Figure 15, which compares translation non-monotonicity shows no aggregate trends. As such, we find that the translation characteristics for E-X language language pairs depends heavily on the individual language pair under consideration.

5.5 X-Y Translation characteristics

Table 9 reports the results of the five measurements for De-Fr and Fr-De translation directions. The results for direct translation pairs are quite different from the X-E and E-X cases, since typically non-English centric translations are done through pivoting. As such, the trends for measurements over Fluency (F), Unaligned Source Words (USW), Unaligned Translation Words (UTW) and Translation Non-Monotonicity (NM) do not show any conclusive evidence of greater paraphrasticity of GPT translations. However, GPT translations still produce greater number of punctuation insertions than the MS Translator system.

5.6 Parallel Data Bias Artifacts

To illustrate the parallel data bias, we analyze the translations on low-quality inputs. The intuition behind our experiment is that low-quality inputs are more likely to correspond to the noisy parts of the parallel data that underlie NMT systems trained on large-scale datasets mined from the web. As such, GPT should outperform NMT systems on such low-quality inputs.

Experiment: We split the test sets into 3 buckets based on perplexity of the source sentence. We notice that the highest perplexity inputs typically correspond to ill-formatted texts, with many inputs pertaining to the e-commerce domain. Such inputs are more likely to *resonate* with the noisy parts of the parallel corpora on which NMT models are typically trained. For example, such high-perplexity inputs might correspond to ill-formed texts scraped from e-commerce websites usually present in parallel corpora. Since we use GPT-2 for obtaining perplexities, we conduct this experiment for E-X language pairs only.

Results: Table 10 presents the results of the experiment across different language pairs. The measurement reported is the average difference in the quality between GPT and MS Translator as measured using COMET-KIWI. We observe that on English to Chinese, English to Japanese and English to Russian language pairs, on which parallel data mining is typically harder owing to a change in script, GPT translations obtain higher performance than MS Translator on the highest perplexity bucket. For low-resource language pairs, GPT gains proportionally in even the lowest perplexity buckets.

Overall, we find that in four out of five high resource language pairs in Table 10, GPT obtains higher improvements in the bucket with the highest input perplexity, when compared to the other lower-perplexity buckets. In the cases of English (Latin script) to Chinese, English to Japanese and English to Russian (Cyrillic script), the differences follow a monotonic order with respect to input perplexity. The results suggest that for these language pairs, GPT does obtain better performance on the lower quality inputs. We attribute this behavior to the parallel data bias. Such parallel data noise biases are likely to be correlated with input domains as well, but we leave such an exploration to future work.

Split	En-De	En-Ru	En-Cs	En-Zh	En-Ja	En-Is	En-Ha
Lowest	-0.02	-0.68	-0.43	0.22	-0.04	-6.34	1.63
Medium	-0.48	-0.41	-0.97	0.47	-0.18	-5.71	1.21
Highest	-0.18	-0.14	-1.10	1.63	0.38	-6.15	1.11

Table 10: **Exploring Parallel Data Bias:** On English to Chinese, Japanese and Russian language pairs, GPT translations obtain higher performance than MS Translator in the highest perplexity bucket. For low-resource language pairs, GPT gains proportionally in even the lowest perplexity buckets.

5.7 Summary

We demonstrated that the computational mechanisms operating behind LLMs and NMT models produce translation artifacts that can be quantitatively differentiated. In this subsection, we summarize our comprehensive characterization of the translations produced by GPT.

Improvements Produced by GPT: For X-E translations, the translations produced by GPT are more fluent, obtaining consistently lower perplexity (as shown in Figure 7). At the same time, GPT translations for X-E language pairs generally incur higher number of unaligned source words (Figure 9) and in general, similar number of unaligned target words (Figure 10). GPT translations are also more non-monotonic, producing translations that involve longer range reorderings (Figure 11). The combination of these results yields an interesting conclusion: that X-E translations by GPT are more fluent and more paraphrastic than the NMT system under investigation (MS Translator), while being faithful to the source. The greater paraphrasticity is not accompanied by content that is unsupported by the source, i.e., the problem of inserted factual content is not a prominent issue in these language pairs.

For E-X translations, GPT incurs a greater number of unaligned source words (USW, Figure 14), along with greater translation non-monotonicity in general (NM, Figure 15), suggesting greater paraphrasticity. However, at the same time, GPT translations incur a slightly higher number of unaligned translation words as well. This suggests that greater paraphrasticity is not the only cause behind the higher USW and NM measurements, and a *less adequate* translation than the NMT system under investigation is a plausible cause behind these observations as well. This is corroborated

by the lower quality measurements for E-X GPT translations obtained previously. In general, we find that for deriving conclusions about GPT translation quality for E-X, it is more important to focus on the single language pair under consideration, i.e. the target non-English language is of critical importance and the effects of language modeling bias cannot be generalized as in the case of X-E translations.

Areas of Improvements: One artifact of the language modeling bias is that GPT inserts end of sentence markers not present in the source with a far greater frequency than the NMT system under investigation. This holds true across both X-E, E-X and X-Y translation directions. Such a proclivity towards greater fluency might not be appropriate for domains wherein a very literal (and faithful) translation is desired. Similarly, greater paraphrasticity might not be appropriate for certain domains. Also, a related area of improvement for future evaluations would be to conduct separate evaluations of fluency and adequacy, in addition to joint adequacy and fluency quality estimation done presently. Instituting a norm of using multi-dimensional automatic quality measurements (Raunak et al., 2022) can provide very targeted signals on differentiating aspects of translation quality, useful especially when there are competing state-of-the-art approaches.

Areas of Application: Our results also suggest that the greater paraphrastic nature of GPT translations could have applications in improving NMT models on the translation of figurative text. Similarly, the greater gains obtained by GPT translations in the highest perplexity buckets of multiple E-X the test sets suggest that GPT translations might be favored over NMT models when the input domain is likely to contain noisy, ill-formed sentences. These two application areas, based on the demonstrated characteristics of GPT translations, offer two avenues that could benefit from a combination of NMT models with GPT. For example, based on the results in Table 10, a hybrid English-Japanese system that could improve upon both MS Translator and GPT translations would be the one wherein the highest perplexity inputs are routed to be translated by GPT whereas the lower perplexity inputs are translated through NMT models (e.g., MS Translator). Such a composition might be able to leverage *complementary* strengths of NMT and LLM systems for translation.

6 Multilingual Capabilities beyond Translation

In this section, we investigate the multilingual capabilities of GPT models beyond translation. Specifically, we aim to assess how well the models perform on emerging reasoning tasks¹⁵ for various languages compared to English. We are interested in understanding the degree of multilingual support that GPT models can offer given their translation performance. That is, can we use the translation performance as a proxy for the multilingual performance on other tasks?

We use MGSM Benchmark (Shi et al., 2022) which is a Multilingual Grade School Math (MGSM) arithmetic reasoning benchmark. The multilingual problems are human translated from the English dataset GSM8K which is English-language human-annotated grade-school math problem dataset. The dataset supports a set of ten languages other than English (EN): Bengali (BN), Chinese (ZH), French (FR), German (DE), Japanese (JA), Russian (RU), Spanish (ES), Swahili (SW), Telugu (TE), and Thai (TH).

Table 11 presents the results on the MSGM benchmark. We first use Native-CoT, which uses prompts and CoT in the native language of each dataset. We observe that *text-davinci-003* surpasses *text-davinci-002* for all languages, highlighting the effectiveness of *text-davinci-003* on multilingual tasks. The performance is especially high on EN, DE, FR and ES, while RU, JA and ZH exhibit lower scores than the Latin languages. The low-resource languages, however, achieve limited performance, indicating the need for better approaches to attain truly multilingual support.

We then use Translate-EN, which translates all prompts and CoT into English. We find that this setup enhances the performance on the non-Latin group (RU, JA and ZH) as well as the low-resource group (TH, TE, BN and SW), although the enhancements are not uniform across languages. Surprisingly, this setup shows a deterioration on the Latin languages.

Our third and final setup is Translate-EN+, which is similar to Translate-EN, but keeps the template in English for all sentences instead of translating it. Stabilizing the template improved results significantly in some languages such as French, Spanish and Russian, and gave comparable scores to Translate-EN in others.

¹⁵as recently studied in the chain of thought paradigm

Setup	EN	DE	FR	ES	RU	ZH	JA	TH	TE	BN	SW
GPT text-davinci-002											
Native-CoT	53.6	36.0	37.6	40.4	28.4	40.0	26.0	10.8	0.4	6.4	11.2
Translate-EN	53.6	46.4	46.4	51.6	48.8	47.2	44.8	41.2	42.8	41.2	37.6
GPT text-davinci-003											
Native-CoT	56.8	54.8	53.6	59.2	40.8	44	38.8	25.9	7.0	14.1	15.2
Translate-EN	56.8	51.6	50.5	49.2	54	48.4	46.7	31.2	44.8	47.6	45.76
Translate-EN+	56.8	49.6	53.2	58.8	57.6	49.2	46.8	30.8	42.5	48.1	44.9
PaLM-540B											
Native-CoT	62.4	49.2	46.4	56.8	48.4	46.8	40.0	52.8	45.6	46.0	35.2
Translate-EN	62.4	57.2	55.2	60.0	59.6	55.6	50.0	50.8	49.6	53.2	51.2

Table 11: GPT performance on MGSM dataset. PaLM-540 results from(Shi et al., 2022).

We observe that, despite the high performance of *text-davinci-003* on the translation of RU, JA and ZH, the performance on MSGM is only moderate. We hypothesize that this may be due to the fact that reasoning tasks benefit greatly from training on programming languages, which is better represented in the top Latin languages, especially with the low proportion of multilingual data in the training data. In contrast, PaLM-540B results from (Shi et al., 2022) show a higher performance with the Native-CoT setup. We hypothesize that this is due to the large multilingual data proportion in its training data, which is 78% English and 22% for other languages (Shi et al., 2022), while GPT data proportion is only 7% non-English (Brown et al., 2020).

These results suggest that translation capability may not be sufficient for the model to exhibit more advanced multilingual reasoning capability, as shown by the poor performance on RU, ZH and JA. We hypothesize that the models acquire their reasoning capabilities through training on natural language multilingual data along with programming languages data, which may limit such capabilities for non-Latin and less represented languages. We think this area warrants more attention from the model developers to provide truly multilingual capabilities across a range of languages.

7 Conclusions and Future Directions

This work presents a comprehensive and in-depth study of the machine translation capabilities of the latest GPT models. Our investigation covers 18 language pairs across four different domains, enabling a broad understanding of the models’ general performance. We also conducted the multilingual reasoning task to examine the interaction between multilinguality and the emergent reasoning capabilities in GPT models. To provide a thorough evaluation

of the models, we employed both human evaluations and the latest neural network-based automatic evaluation metrics together with the conventional machine translation evaluation metrics. In addition, we conducted extensive analysis, providing an in-depth examination of various phenomena in GPT models’ translation outputs and their comparisons to state-of-the-art NMT systems.

As a result, our findings demonstrate that GPT systems can produce highly fluent and competitive translation outputs even in the zero-shot setting especially for the high-resource language translations. By utilizing the in-context learning capability of GPT models with few-shot examples, we were able to further improve translation quality. Additionally, we demonstrated that a hybrid approach, combining the latest NMT systems with GPT models, can achieve state-of-the-art translation quality.

While the use of LLMs in machine translation is a rapidly developing area, there are many research directions that can be explored to improve the quality and understanding of machine translation. Below are some of the important areas that we focus on:

- **Underrepresented languages:** Our study has shown that GPT models, still struggle with underrepresented languages, which makes it a critical research question to explore how to improve the translation quality for these languages.
- **In-context learning:** GPT models have shown great potential for in-context learning, which can be leveraged to generate different styles or nuanced translations. Future research can explore how to better utilize this capability to improve translation quality.
- **Model fusion:** The use of large-scale LLMs like GPT can be computationally expensive,

and therefore, exploring how to more efficiently utilize them is an important research question. We are investigating more sophisticated fusion techniques that can achieve higher quality and efficiency.

- **Better metrics:** The limitations of lexical matching metrics can mislead translation quality assessment. Therefore, developing metrics that can measure the contextual correctness of LLM-generated translations is essential. Future research can also explore new ways to evaluate the quality of machine translations more accurately, especially when using LLMs.

Overall, our study provides valuable insights into the strengths and weaknesses of GPT models for machine translation and opens up opportunities for future improvements and developments in this field. We investigated how GPT models can transform machine translation as it is doing with other generative tasks. We demonstrated that these models excel at translating well-represented languages in their training data, but they face challenges with less-resourced languages. We also assessed translation and reasoning tasks and detected discrepancies in the level of support of the tasks even for the same languages. One of the main benefits of training such costly models is to achieve high performance across diverse tasks and languages, but this demands more data across languages. Which poses several challenges for models scalability, diversity, and fairness. As a future research direction, we propose to tackle the challenge of enabling truly multilingual capability for such models that would enable the same capabilities across languages.

Limitations

We conducted our evaluation on 18 translation directions with reliable test sets and baselines. The conclusion of the study should be taken in this context and not generalized to other languages without further evaluation. While a more comprehensive evaluation is needed on more languages, we should be cautious about drawing conclusions from low quality testsets or weaker baselines which are usually dominating the research results for low resource languages.

One of the limitations of this study is the inadequacy of current automatic evaluation metrics to capture the quality of GPT outputs accurately. We

found that lexical comparison based metrics such as BLEU or chrF gave misleading signals, and that document-level evaluation had limited capability to realize the effect of context-based translation. These limitations stem from the inherent challenges of evaluating natural language generation systems, especially for complex tasks such as machine translation. Therefore, we complemented our automatic evaluation with a comprehensive analysis, considering all metrics together, as well as human evaluation and qualitative analysis to cover a broad range of phenomena. We recommend that readers consider the overall evaluations as a whole, rather than relying solely on a specific metric, to better understand the quality of GPT models' machine translation capabilities.

Ethics Statement

In our study, we evaluated the quality of GPT translations using various quantitative metrics. However, we acknowledge that these models may harbor language-specific biases and produce translations that perpetuate stereotypes and misinformation. One form of bias that we have identified is that the models perform better for some languages than others, as shown in §3.3 and §3.4. This may create unfairness and unequal outcomes for users paying the same cost. Furthermore, the models may amplify stereotypes from the training data, leading to inaccurate translations. For instance, the model may fail to correctly translate a sentence with a female name like “Julia” as a male. Lastly, we also observed instances of false insertions in translations (misinformation) and hallucinations, especially in E-X translation directions. We are committed to addressing these issues and mitigating biases and misinformation in our research and future work.

References

- Sweta Agrawal, Chunting Zhou, Mike Lewis, Luke Zettlemoyer, and Marjan Ghazvininejad. 2022. [In-context examples selection for machine translation](#).
- Farhad Akhbardeh, Arkady Arkhangorodsky, Magdalena Biesialska, Ondřej Bojar, Rajen Chatterjee, Vishrav Chaudhary, Marta R Costa-jussà, Cristina España-Bonet, Angela Fan, Christian Federmann, et al. 2021. [Findings of the 2021 conference on machine translation \(WMT21\)](#). In *Proceedings of the Sixth Conference on Machine Translation*, pages 1–88.

- Anonymous. 2023a. [Dissecting in-context learning of translations in gpt-3](#). Anonymous preprint under review.
- Anonymous. 2023b. [Does gpt-3 produces less literal translations?](#) Anonymous preprint under review.
- Loic Barrault, Ondrej Bojar, Fethi Bougares, Rajen Chatterjee, Marta R. Costa-jussa, Christian Federmann, Mark Fishel, Alexander Fraser, Markus Freitag, Yvette Graham, Roman Grundkiewicz, Paco Guzman, Barry Haddow, Matthias Huck, Antonio Jimeno Yepes, Philipp Koehn, Tom Kocmi, Andre Martins, Makoto Morishita, and Christof Monz, editors. 2021. [Proceedings of the Sixth Conference on Machine Translation](#). Association for Computational Linguistics, Online.
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. [Language models are few-shot learners](#). *CoRR*, abs/2005.14165.
- Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, et al. 2022. [Palm: Scaling language modeling with pathways](#). *arXiv preprint arXiv:2204.02311*.
- Marta R Costa-jussà, James Cross, Onur Çelebi, Maha Elbayad, Kenneth Heafield, Kevin Heffernan, Elahe Kalbassi, Janice Lam, Daniel Licht, Jean Maillard, et al. 2022. [No language left behind: Scaling human-centered machine translation](#). *arXiv preprint arXiv:2207.04672*.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. [BERT: Pre-training of deep bidirectional transformers for language understanding](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics.
- Zi-Yi Dou and Graham Neubig. 2021. [Word alignment by fine-tuning embeddings on parallel corpora](#). In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 2112–2128, Online. Association for Computational Linguistics.
- Angela Fan, Shruti Bhosale, Holger Schwenk, Zhiyi Ma, Ahmed El-Kishky, Siddharth Goyal, Mandeep Baines, Onur Celebi, Guillaume Wenzek, Vishrav Chaudhary, et al. 2021. [Beyond english-centric multilingual machine translation](#). *The Journal of Machine Learning Research*, 22(1):4839–4886.
- Fangxiaoyu Feng, Yinfei Yang, Daniel Cer, Naveen Arivazhagan, and Wei Wang. 2020. [Language-agnostic bert sentence embedding](#).
- Markus Freitag, Ricardo Rei, Nitika Mathur, Chiklu Lo, Craig Stewart, Eleftherios Avramidis, Tom Kocmi, George Foster, Alon Lavie, and André Martins. 2022. [Results of wmt22 metrics shared task: Stop using bleu - neural metrics are better and more robust](#). In *Proceedings of the Seventh Conference on Machine Translation*, pages 46–68, Abu Dhabi.
- Xavier Garcia, Yamini Bansal, Colin Cherry, George Foster, Maxim Krikun, Fangxiaoyu Feng, Melvin Johnson, and Orhan Firat. 2023. [The unreasonable effectiveness of few-shot learning for machine translation](#). *arXiv preprint arXiv:2302.01398*.
- Tanya Goyal, Junyi Jessy Li, and Greg Durrett. 2022. [News summarization and evaluation in the era of gpt-3](#). *arXiv preprint arXiv:2209.12356*.
- Young Jin Kim, Ammar Ahmad Awan, Alexandre Muzio, Andres Felipe Cruz Salinas, Liyang Lu, Amr Hendy, Samyam Rajbhandari, Yuxiong He, and Hany Hassan Awadalla. 2021. [Scalable and efficient moe training for multitask multilingual models](#). *arXiv preprint arXiv:2109.10465*.
- Tom Kocmi, Rachel Bawden, Ondřej Bojar, Anton Dvorkovich, Christian Federmann, Mark Fishel, Thamme Gowda, Yvette Graham, Roman Grundkiewicz, Barry Haddow, et al. 2022a. [Findings of the 2022 conference on machine translation \(wmt22\)](#). In *Proceedings of the Seventh Conference on Machine Translation (WMT)*, pages 1–45.
- Tom Kocmi, Rachel Bawden, Ondřej Bojar, Anton Dvorkovich, Christian Federmann, Mark Fishel, Thamme Gowda, Yvette Graham, Roman Grundkiewicz, Barry Haddow, Rebecca Knowles, Philipp Koehn, Christof Monz, Makoto Morishita, Masaaki Nagata, Toshiaki Nakazawa, Michal Novák, Martin Popel, Maja Popović, and Mariya Shmatova. 2022b. [Findings of the 2022 conference on machine translation \(wmt22\)](#). In *Proceedings of the Seventh Conference on Machine Translation*, pages 1–45, Abu Dhabi. Association for Computational Linguistics.
- Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. 2022. [Large language models are zero-shot reasoners](#).
- Yinhan Liu, Jiatao Gu, Naman Goyal, Xian Li, Sergey Edunov, Marjan Ghazvininejad, Mike Lewis, and Luke Zettlemoyer. 2020. [Multilingual denoising pre-training for neural machine translation](#). *Transactions of the Association for Computational Linguistics*, 8:726–742.

- Sewon Min, Xinxu Lyu, Ari Holtzman, Mikel Artetxe, Mike Lewis, Hannaneh Hajishirzi, and Luke Zettlemoyer. 2022. [Rethinking the role of demonstrations: What makes in-context learning work?](#) In *EMNLP*.
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. 2022. [Training language models to follow instructions with human feedback](#).
- Maja Popović. 2015. [chrF: character n-gram F-score for automatic MT evaluation](#). In *Proceedings of the Tenth Workshop on Statistical Machine Translation*, pages 392–395, Lisbon, Portugal. Association for Computational Linguistics.
- Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. 2019. [Language models are unsupervised multitask learners](#). *OpenAI blog*, 1(8):9.
- Vikas Raunak and Arul Menezes. 2022. [Finding memo: Extractive memorization in constrained sequence generation tasks](#). In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 5153–5162, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.
- Vikas Raunak, Arul Menezes, and Marcin Junczys-Dowmunt. 2021. [The curious case of hallucinations in neural machine translation](#). In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 1172–1183, Online. Association for Computational Linguistics.
- Vikas Raunak, Matt Post, and Arul Menezes. 2022. [SALTED: A framework for SALient long-tail translation error detection](#). In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 5163–5179, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.
- Ricardo Rei, Jos   G. C. de Souza, Duarte Alves, Chrysoula Zerva, Ana C Farinha, Taisiya Glushkova, Alon Lavie, Luisa Coheur, and Andr   F. T. Martins. 2022a. [Comet-22: Unbabel-ist 2022 submission for the metrics shared task](#). In *Proceedings of the Seventh Conference on Machine Translation*, pages 578–585, Abu Dhabi. Association for Computational Linguistics.
- Ricardo Rei, Marcos Treviso, Nuno M. Guerreiro, Chrysoula Zerva, Ana C. Farinha, Christine Maroti, Jos   G. C. de Souza, Taisiya Glushkova, Duarte M. Alves, Alon Lavie, Luisa Coheur, and Andr   F. T. Martins. 2022b. [Cometkiwi: Ist-unbabel 2022 submission for the quality estimation shared task](#).
- Moritz Schaeffer and Michael Carl. 2014. [Measuring the cognitive effort of literal translation processes](#). In *Proceedings of the EACL 2014 Workshop on Humans and Computer-assisted Translation*, pages 29–37, Gothenburg, Sweden. Association for Computational Linguistics.
- Andrea Schioppa, David Vilar, Artem Sokolov, and Katja Filippova. 2021. [Controlling machine translation for multiple attributes with additive interventions](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 6676–6696, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Freda Shi, Mirac Suzgun, Markus Freitag, Xuezhi Wang, Suraj Srivats, Sorous Vosoughi, Hyung Won Chung, Yi Tay, Sebastian Ruder, Denny Zhou, Dipanjan Das, and Jason Wei. 2022. [Language models are multilingual chain-of-thought reasoners](#).
- Zewei Sun, Mingxuan Wang, Hao Zhou, Chengqi Zhao, Shujian Huang, Jiajun Chen, and Lei Li. 2020. [Rethinking document-level neural machine translation](#).
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. 2017. [Attention is all you need](#). *CoRR*, abs/1706.03762.
- David Vilar, Markus Freitag, Colin Cherry, Jiaming Luo, Viresh Ratnakar, and George Foster. 2022. [Prompting palm for translation: Assessing strategies and performance](#).
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Ed Chi, Quoc Le, and Denny Zhou. 2022. [Chain of thought prompting elicits reasoning in large language models](#). *arXiv preprint arXiv:2201.11903*.
- Biao Zhang, Barry Haddow, and Alexandra Birch. 2023. [Prompting large language model for machine translation: A case study](#).
- Mike Zhang and Antonio Toral. 2019. [The effect of translationese in machine translation test sets](#). In *Proceedings of the Fourth Conference on Machine Translation (Volume 1: Research Papers)*, pages 73–81, Florence, Italy. Association for Computational Linguistics.

A Prompt Templates

Translate this into 1. [target language]:

[shot n source]

1. [shot n reference]

Translate this into 1. [target language]:

[input]

1.

Figure 16: Prompt template for sentence-level translation. We use the same instruction and format as recommended in the OpenAI playground for the default sentence-level translation task ¹⁶. Keeping the prompt format the same allows us to potentially leverage the benefits of the underlying instruction finetuning protocol to the full extent.

Translate this sentence from [source language] to [target language], Source:

[source sentence]

Target:

Figure 17: Zero-shot prompt template for ChatGPT.

Document:
[shot 1 source]

[shot 2 source]

[shot n source]

Translate each line in document into [target language].

Translated Document:

[shot 1 reference]

[shot 2 reference]

[shot n reference]

Document:
[sentence 1 from context window]

[sentence n from context window]

Translate each line in document into [target language].

Translated Document:

Figure 18: Prompt template for document translation.

B Few-shot Example Selection Data Pool

Language	# of sentences	
	Raw	Cleaned
CS-EN	193.5M	175.5M
EN-CS	167.9M	151.5M
DE-EN/EN-DE	295.8M	289.1M
IS-EN/EN-IS	4.4M	3.7M
JA-EN/EN-JA	33.9M	33.1M
ZH-EN	55.2M	50.4M
EN-ZH	35.5M	31.2M
UK-EN/EN-UK	50.6M	45.5M
RU-EN/EN-RU	75.0M	65.7M
HA-EN/EN-HA	7.2M	727K
FR-DE/DE-FR	17.7M	15.6M

Table 12: Size of the data pool for the few-shot example selections for each translation direction. Raw column shows the size of the original dataset which is from the WMT training dataset and Cleaned column shows high quality data after the cleaning from the original dataset.

C Human Evaluation details

LP	WMT-HE	GPT-HE	Delta-HE	WMT-COMETkiwi	GPT-COMETkiwi	Delta COMETkiwi
DE-EN	93.8	92.0	-1.8	81.4	81.4	0.0
CS-EN	84.1	85.7	1.7	82.5	82.5	0.0
JP-EN	76.5	74.7	-1.8	80.3	80.8	0.5
ZH-EN	77.4	80.0	2.5	77.7	78.5	0.8
RU-EN	88.0	87.2	-0.8	81.7	82.9	1.2
UK-EN	83.3	80.1	-3.1	81.5	80.3	-1.2
IS-EN	91.4	86.8	-4.6	81.4	80.2	-1.2
HA-EN	84.8	75.9	-8.9	74.5	68.5	-6.0
EN-DE	95.0	93.4	-1.5	83.6	82.9	-0.7
EN-CS	85.9	80.5	-5.4	84.2	83.3	-0.9
EN-JP	79.3	76.0	-3.3	85.8	85.3	-0.5
EN-ZH	81.6	81.9	0.4	82.0	82.0	0.0
EN-RU	88.6	83.7	-4.9	84.4	82.2	-2.2
EN-UK	85.2	76.9	-8.3	83.4	80.6	-2.8
EN-IS	92.1	69.2	-22.9	81.8	74.1	-7.7
EN-HA	80.2	55.5	-24.7	61.5	58.5	-3.0
FR-DE	85.7	84.6	-1.1	80.7	80.2	-0.5
DE-FR	85.0	86.5	1.5	79.5	80.7	1.2

Table 13: Human Evaluation and COMETkiwi Results of WMT Best Systems and GPT in details.

C.1 Human Evaluation Analysis

As shown in the examples for the first characteristic in Table 14, GPT is good at handling misspelled words or unclosed quotes and ends translations without missing any semantic information. Also, GPT translations are reasonable for partial or incomplete colloquial source sentences whereas WMT-Best translations have extra or missing content (the third characteristic).

On the other hand, GPT tends to generate unnatural translations for sentences with unusual or complex expressions (the second characteristic). GPT also generates hallucinated words in its translations as in the fourth characteristic. For instance, the first GPT example translation uses “udon” (“noodle” in English) although the source never mentions it. In the second GPT example translation, the word “PM” (Prime Minister) is translated into a non-existing Japanese word that sound like an existing jargon. Unlike oscillatory hallucinations (i.e., repeating the same words or phrases many times) (Raunak et al., 2021), these subtle, word-level hallucinations are problematic as they are hard to be detected and may not severely affect metric scores (Raunak et al., 2022).

Characteristic	Examples
	<p>SRC: please try to <u>perfrm thsi procedures</u>”</p> <p>GPT: この手順を試してください。(84)</p> <p>WMT: <u>perfrm thsi procedures</u>” を実行してみてください。(53)</p>
GPT produces reasonable translations for erroneous, short or unconventional source sentences (EN-JP)	<p>ENU: Ms Truss <u>said</u>: ”The UK is providing vital humanitarian assistance in Afghanistan this winter.</p> <p>GPT: Truss 氏は、「今冬、イギリスはアフガニスタンで重要な人道支援を提供しています」と述べました。(85)</p> <p>WMT: トラス女史は、「英国はこの冬、アフガニスタンで重要な人道支援を行っている。(68)</p>
	<p>ENU: Or they’ll claim that <u>shredded ballots</u> were <u>fed to chickens</u> that were then incinerated.</p> <p>GPT: または、彼らは、<u>ちぎれた投票用紙が鶏に与えられ</u>、その後焼却されたと主張します。(50)</p> <p>WMT: または、細断された投票用紙が鶏の餌となり、その後焼却されたと主張するだろう。(85)</p>
GPT produces undesirable direct / literal translations (EN-JP)	<p>ENU: The Trump supporter has also pushed for <u>debunked claims of voter fraud in Arizona</u>.</p> <p>GPT: トランプ支持者は、アリゾナの<u>投票不正の破棄された主張</u>も推し進めています。(61)</p> <p>WMT: トランプ大統領の支持者は、アリゾナ州での不正投票疑惑の暴露も推進している。(90)</p>
	<p>JPN: ある人気役者さんが、言っていました。</p> <p>GPT: A popular actor said, (52)</p> <p>WMT: A popular actor once said, ”It’s not like that. (37)</p>
GPT produces reasonable translations for short, colloquial source sentences (JP-EN)	<p>JPN: …っていう話をどっかでちらっと見たんですが、どこでみたのか思い出せません…</p> <p>GPT: I saw something about it somewhere, but I can’t remember where... (67)</p> <p>WMT: But I can’t remember where I saw it... (42)</p>
	<p>JPN: めっちゃよく見たらすこしわかるかもしれないけど<u>うなぎ</u>は一切使ってないえせうな丼。</p> <p>GPT: It might look like it has some <u>udon</u> in it, but there’s actually no <u>udon</u> in this bowl. (13)</p> <p>WMT: If you look very closely, you may understand a little, but they don’t use any eel at all. (56)</p>
GPT generates hallucinated words (JP-EN, EN-JP)	<p>ENU: Greece’s Defense Ministry also confirmed that the deal with Paris was ”final,” since it had been negotiated at the ”highest level possible” and ”personally announced” by the Greek <u>PM</u> Kyriakos Mitsotakis.</p> <p>GPT: ギリシャの農林水産省は、パリとの取引が「最終的なもの」であることを確認しました。これは、ギリシャの<u>プレミアミンストラ</u> Kyriakos Mitsotakis が「最高レベルで交渉」し、「個人的に発表」したためです。(42)</p> <p>WMT: ギリシャ国防省はまた、パリとの合意は「可能な限り最高レベル」で交渉され、ギリシャのキリアコス・ミツォタキス首相によって「個人的に発表された」ため、「最終的」であることを確認した。(92)</p>

Table 14: Qualitative analysis for GPT English from/to Japanese translations. These examples show some positive and negative characteristics indicated by human item scores (numbers in parentheses).